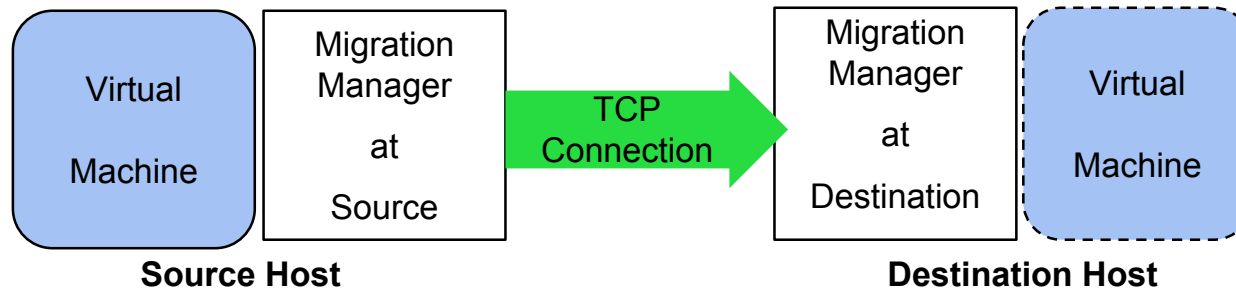


Live Migration of Virtual Machines

Pre-copy :Christopher Clarke, Keir Fraser, et. al. NSDI 2005

Post-copy: Hines, Deshpande, Gopalan, VEE 2009

What is live VM migration?



- Move a VM from one physical machine to another even as its applications continue to execute during migration
- Live VM migration usually involves
 - Migrating memory state
 - Migrating CPU state
 - Optionally, migrating virtual disk state
- Migration managers at source and destination
 - Connect via TCP connection
 - At source, the migration manager maps the guest VM's memory and execution state
 - Transfers VM's pages to the target migration manager over TCP connection.
 - At destination, the migration manager restores the VM's state and resumes execution
 - Migration manager examples: xend for Xen, QEMU for KVM

Why Live VM Migration?

- Why Migrate?

- Load Balancing
 - Move VMs from highly loaded servers to lightly loaded servers
- Server maintenance
 - When server needs to be upgraded
- Energy savings
 - Move out VMs before shutting down servers to reduce energy usage

- Why live?

- To keep long-running jobs alive
- To keep network connections alive
- Broadly, to avoid disruptions to users of VM

- Why VM?

- Why not migrate individual processes?
- Process migration may leave residual dependencies (state) at source host
 - E.g. system call redirection, shared memory, open files, inter-process communication, etc.

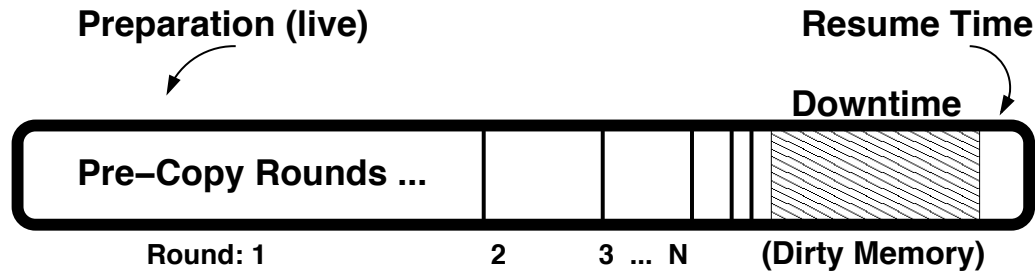
Performance Goals in Live Migration

- Minimizing Downtime
- Reducing total migration time
- Avoiding interference with normal system activity
- Minimizing network activity

Migrating Memory

- Pure stop-and-copy
 - Freeze VM at source,
 - Copy the VM's pseudo-physical memory contents to target,
 - Restart VM at target
 - Long downtime.
 - Minimal total migration time = downtime
- Pure Demand Paging:
 - Freeze VM at source,
 - Copy minimal execution context to target
 - PC, Registers, non-pageable memory
 - Restart VM at target,
 - Pull memory contents from source as and when needed
 - Smaller downtime
 - Slooooooow warm-up phase at target during page-faults across network

Pre-copy migration



(a) Pre-Copy Timeline

----->
Time

- DON'T freeze VM at source
 - Let the VM continue to run
- Copy VM's pseudo-physical memory contents to target over multiple iterations
 - First iteration → copy all pages.
 - Each subsequent iteration → copy pages that were dirtied by the VM during the previous iteration
- Do a short stop-and-copy when number of dirty pages is “small enough”.
- But what if number of dirty pages never converges to a small enough number?
 - After a fixed number of iterations, give up and stop-and-copy.

So what's the catch?

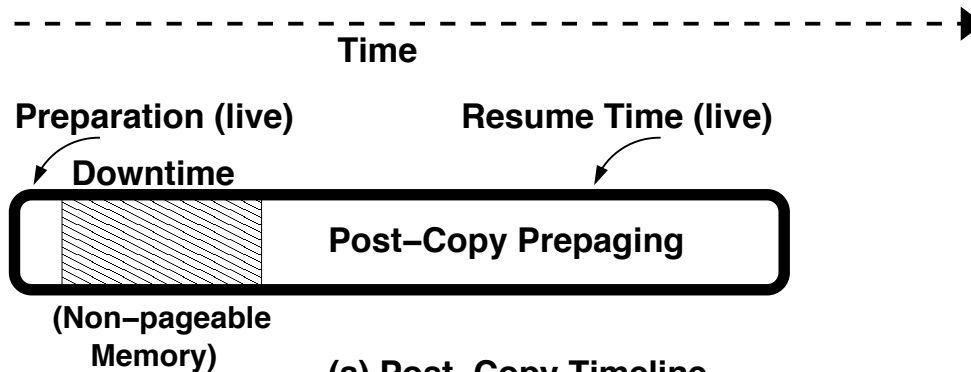
How do we track dirtied pages?

- Mark the VM's memory pages as read-only after each iteration.
- Trap write operations via hypervisor to xend and track dirtied pages.
- Reset after each iteration
- Works well as long as writes are infrequent

Optimizations


- Limit the bandwidth used by migration
 - To minimize impact on running services
- Stun Rogue Processes
 - Those that don't stop dirtying memory
- Free Page Cache Pages
 - Can be re-cached at target
 - Potential performance hit

Post-copy migration



- Freeze the VM first
- Migrate CPU state and minimum state to destination
- Start VM at the target, but without its memory!
- Transfer memory by concurrently doing the following
 - Demand paging over network
 - Actively pushing from source
 - Hopefully most pages will be pushed BEFORE they are demand paged.
- Advantage:
 - Each page transferred over the network only once.
 - Deterministic total migration time
- Disadvantage:
 - Cold start penalty at the destination
 - If migration fails, then VM is lost.

Hybrid pre/post-copy



- Combines the benefits & drawbacks of both
 1. Perform one or more rounds of live pre-copy rounds
 2. Pause VM and transfer execution state
 3. Use post-copy to transfer any remaining dirty pages from source

Migrating Network Connections

- Within a LAN,
 - the migrated VM carries its IP address, MAC address, and all protocol state, including any open sockets
 - Backward (re)learning delay at the network switches
 - Switches need to re-learn the new location of migrated VM's MAC address
 - Solution: Send an unsolicited ARP reply from the target host.
 - Intermediate switches will re-learn automatically.
 - Few in-flight packets might get lost.
- Across a WAN (wide-area network)
 - Source and destination subnets may have different IP addresses.
 - Active network connections may need to be tunneled via VPN or similar mechanisms.

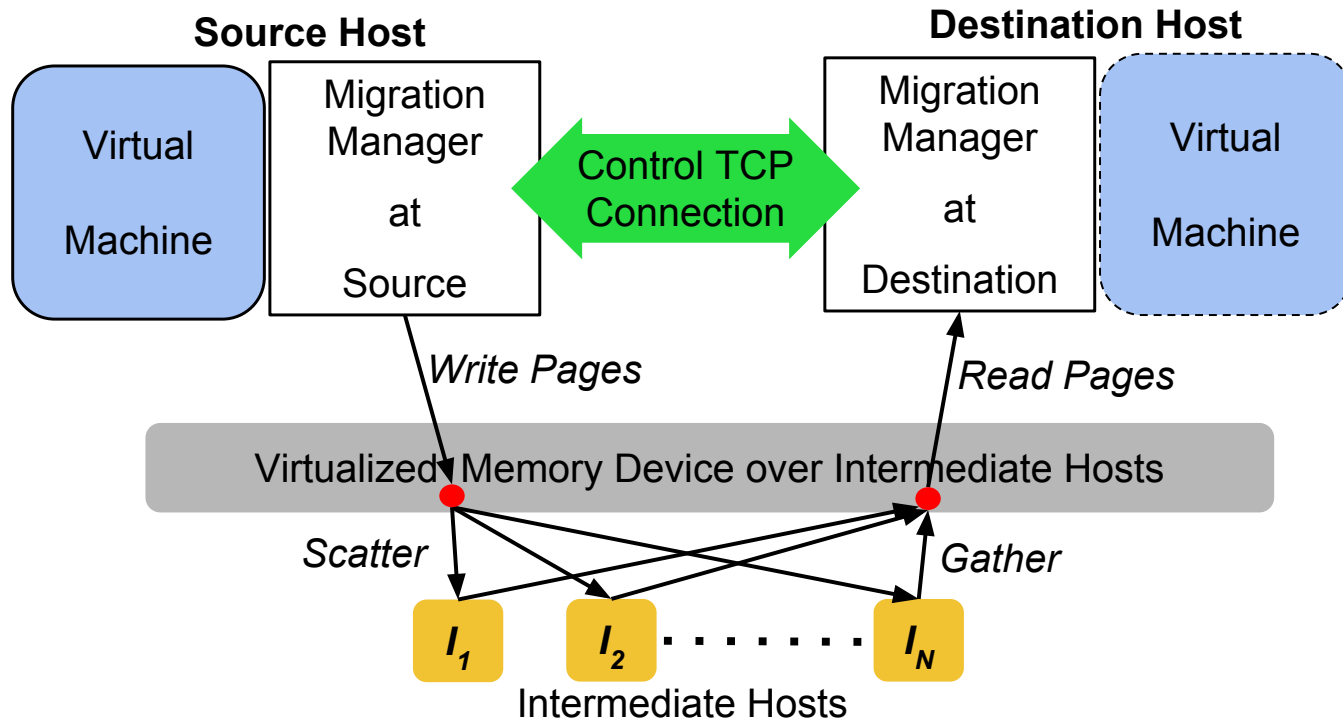
Storage Migration

- Many gigabytes of local disk image possible.
- For LAN
 - Assume the storage is over the network and remains accessible from the new target machine.
 - E.g. Network File System (NFS), or Network Block Device(NBD), or iSCSI etc.
- For WAN
 - Disk image may need to be transferred.
 - Can use pre-copy or post-copy for disk images,
 - Combined bandwidth saving optimizations such as compression, and/or de-duplication.

Self Migration


- Guest OS migrates itself (mostly)
 - No migration manager needed at source!
- Migration stub needed at destination
 - Fresh VM that receives guest state and restores migrated VM's state.
- Challenge:
 - OS must continue to execute while transferring its final state.
 - Perform a careful (complicated) 2-stage checkpoint and copy.

Scatter-Gather migration



Scatter-Gather migration: The VM's state is transferred through intermediaries. A direct connection between the source and destination carries control information, faulted pages, and some actively pushed pages.

Multi-VM (Gang) Migration



- De-duplicate memory pages to reduce network traffic.
- Identify identical pages across multiple VMs
 - By comparing byte-wise (expensive), or checksum (cheaper)
- Send only one copy of identical page to destination node
- Destination Node replicates the pages to multiple VMs.