


System Calls

Kartik Gopalan


System Calls

- Modern CPUs support at least two levels of privileges:
 - User mode - application execute at this level
 - Supervisor mode - OS (kernel) code executes at this level
- System calls
 - Interface to allow User-level processes to safely invoke OS routines for privileged operations.
 - Safely transfer control from lower privilege level (user mode) to higher privilege level (supervisor mode), and back.



System Call table

- Protected entry points into the kernel for each system call
 - We don't want application to randomly jump into any part of the OS code.
- Syscall table is usually implemented as an array of function pointers, where each function implements one system call
- Syscall table is indexed via system call number



Steps in system call execution

User process	Invoke syscall using, say, SYSENTER instruction (arguments in registers/stack)
CPU	Switch CPU to <u>supervisor</u> mode. Jump to entry point in kernel.
Kernel	Save process state Lookup Syscall table. Invoke syscall.
Kernel	Optionally Block process if it needs to wait for I/O or other events. Return process to ready state when woken.
Kernel	Restore saved process state SYSEXIT
CPU	Switch CPU to <u>user</u> mode Return to user process
User Process	Return from system call. Continue

Syscall Usage

- To make it easier to invoke system calls, OS writers normally provide a library that sits between programs and system call interface.
 - Libc, glibc, etc.
- This library provides wrapper routines
- Wrappers hide the low-level details of
 - Preparing arguments
 - Passing arguments to kernel
 - Switching to supervisor mode
 - Fetching and returning results to application.
- Helps to reduce OS dependency and increase portability of programs.

Implementing System Calls

Steps in writing a system call

1. Create an entry for the system call in the kernel's `syscall_table`
 - User processes trapping to the kernel (through `SYS_ENTER` or int 0x80) find the syscall function by indexing into this table.
2. Write the system call code as a kernel function
 - Be careful when reading/writing to user-space
 - Use `copy_to_user()` or `copy_from_user()` routines.
 - These perform sanity checks.
3. Implement a user-level wrapper to invoke your system call
 - Hides the complexity of making a system call from user applications.
 - See *man syscall*

Step 1: Create a sys_call_table entry (for 64-bit x86 machines)

- Syscall table initialized in [arch/x86/entry/syscall_64.c](#)

- [arch/x86/entry/syscalls/syscall_64.tbl](#)

```
#
# 64-bit system call numbers and entry vectors
#
# The format is:
# <number> <abi> <name> <entry point>
#
# The abi is "common", "64" or "x32" for this file.
...
309 common          getcpu                sys_getcpu
310 64              process_vm_readv      sys_process_vm_readv
311 64              process_vm_writev     sys_process_vm_writev
312 common          kcmp                  sys_kcmp
313 common          foo                   sys_foo
```


Step 2: Write the system call handler

- System call with no arguments and integer return value

```
SYSCALL_DEFINE0(foo) {  
    printk (KERN_ALERT "sys_foo: pid is %d\n", current->pid);  
    return current->pid;  
}
```

- Syscall with one primitive argument

```
SYSCALL_DEFINE1(foo, int, arg) {  
    printk (KERN_ALERT "sys_foo: Argument is %d\n", arg);  
    return arg;  
}
```

- To see system log:
 - `dmesg`
 - `less /var/log/kern.log`

Step 2: Write the system call handler

- Verifying argument passed by user space

```
SYSCALL_DEFINE1(close, unsigned int, fd)
{
    struct file * filp;
    struct files_struct *files = current->files;
    struct fdtable *fdt;
    spin_lock(&files->file_lock);
    fdt = files_fdt(files);

    if (fd >= fdt->max_fds)
        goto out_unlock;
    filp = fdt->fd[fd];
    if (!filp)
        goto out_unlock;

    ...
out_unlock:
    spin_unlock(&files->file_lock);
    return -EBADF;
}
```

- Call-by-reference argument
 - User-space pointer sent as argument.
 - Data to be copied back using the pointer.

```
SYSCALL_DEFINE3(read, unsigned int, fd,
                char __user *, buf, size_t, count)
{
    ...

    if( !access_ok( VERIFY_WRITE, buf, count))
        return -EFAULT;

    ...
}
```

Step 3: Invoke syscall handler from user space

- Use the **syscall(...)** library function.
 - Do a "man syscall" for details.
- For instance, for a no-argument system call named foo(), you'll call
 - `ret = syscall(__NR_sys_foo);`
 - Assuming you've defined `__NR_sys_foo` earlier
- For a 1 argument system call named foo(arg), you call
 - `ret = syscall(__NR_sys_foo, arg);`
- and so on for 2, 3, 4 arguments etc.
- For this method, check
 - <https://developer.ibm.com/tutorials/1-system-calls/>

Step 3: Invoke your new handler from user space

```
#include <stdio.h>
#include <errno.h>
#include <unistd.h>
#include <linux/unistd.h>
// define the new syscall number. Standard syscalls are defined in linux/unistd.h
#define __NR_sys_foo 333
int main(void)
{
    int ret;
    while(1) {
        // making the system call
        ret = syscall(__NR_sys_foo);
        printf("ret = %d errno = %d\n", ret, errno);
        sleep(1);
    }
    return 0;
}
```