Memory Management

Kartik Gopalan

References:
- Chapter 3, Modern Operating Systems, Andrew S. Tanenbaum
- https://en.wikipedia.org/wiki/Page_(computer_memory)
- https://en.wikipedia.org/wiki/Page_table
- https://en.wikipedia.org/wiki/Virtual_memory
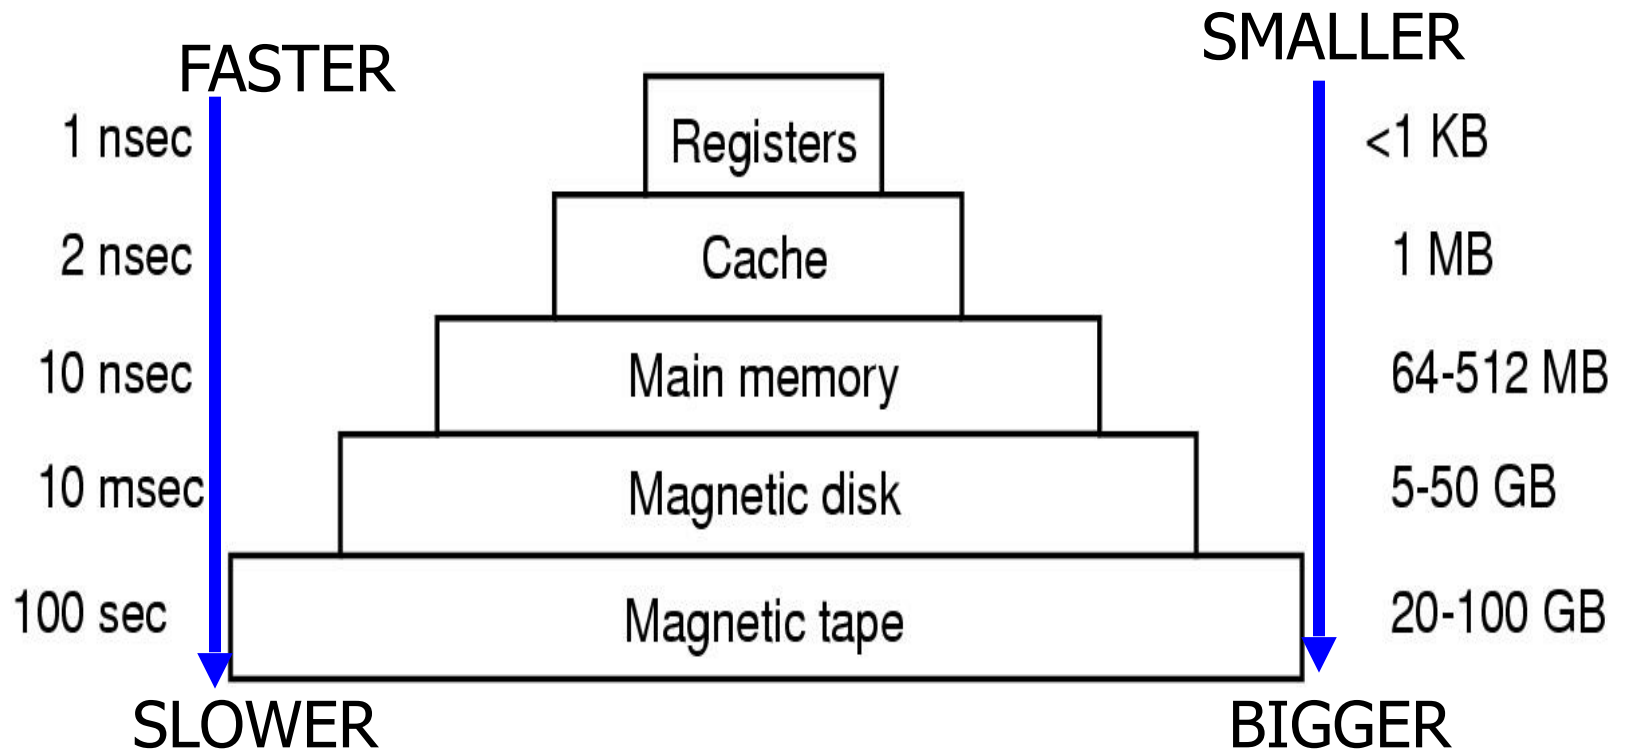
Memory Management

- Ideally programmers want memory that is
  - large
  - fast
  - persistent (non-volatile)

Memory Hierarchy

- Registers & Cache
  - small amount of fast, expensive, volatile memory
- Main memory
  - some medium-speed, medium price, volatile/persistent memory
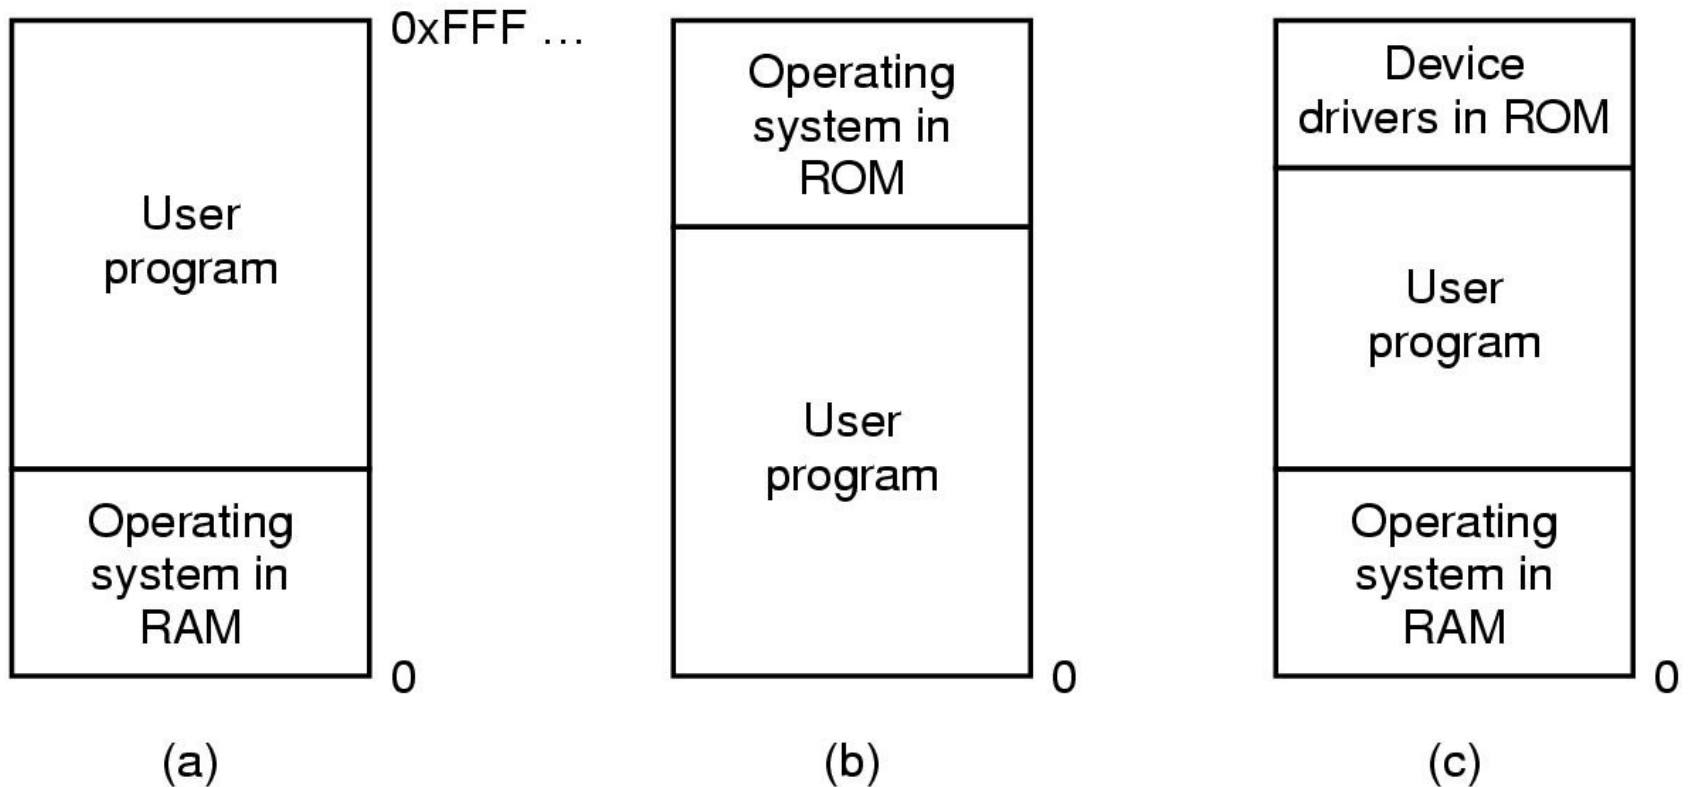- Disk & Tape
  - Lots of slow, cheap, persistent, storage

Typical access time

Typical capacity

FASTER

SMALLER

| Typical access time		Typical capacity
1 nsec	Registers	<1 KB
2 nsec	Cache	1 MB
10 nsec	Main memory	64-512 MB
10 msec	Magnetic disk	5-50 GB
100 sec	Magnetic tape	20-100 GB

SLOWER

BIGGER

3

Basic Memory Management
"Mono-programming" without Swapping or Paging

0xFFF …

(a)
- User program
- Operating system in RAM
- 0

(b)
- Operating system in ROM
- User program
- 0

(c)
- Device drivers in ROM
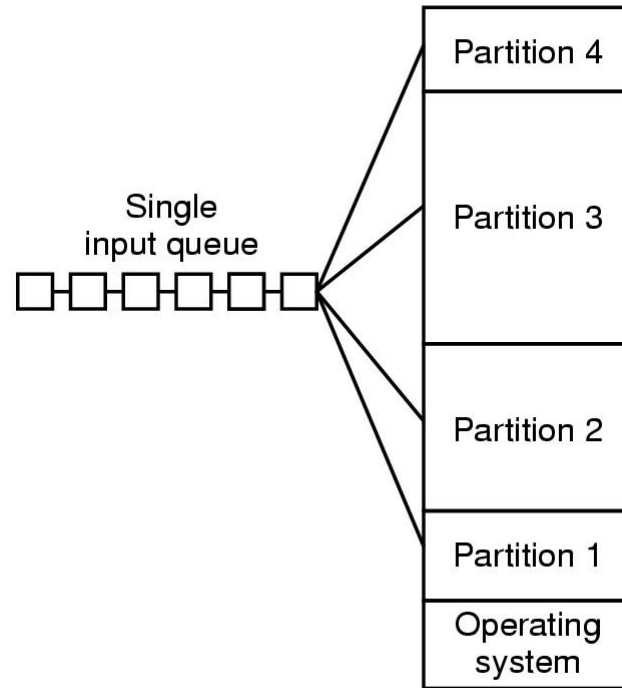- User program
- Operating system in RAM
- 0

Three simple ways of organizing memory
- an operating system with one user process
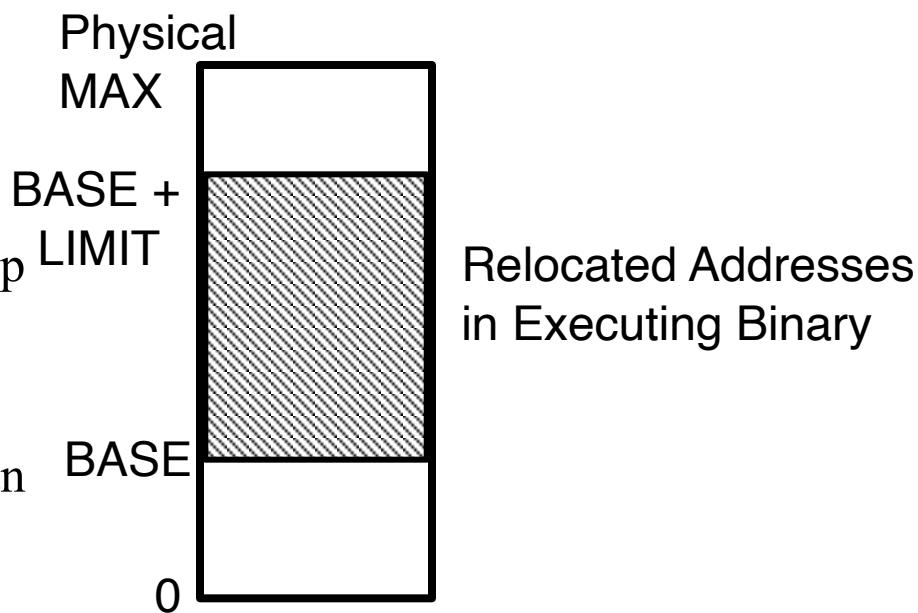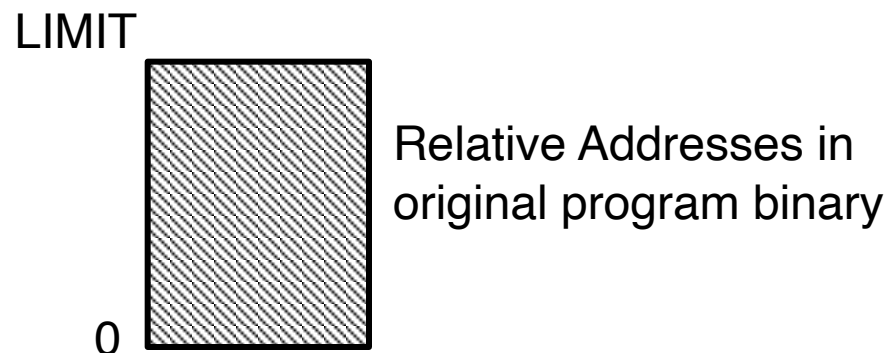
Multiprogramming with Fixed Partitions
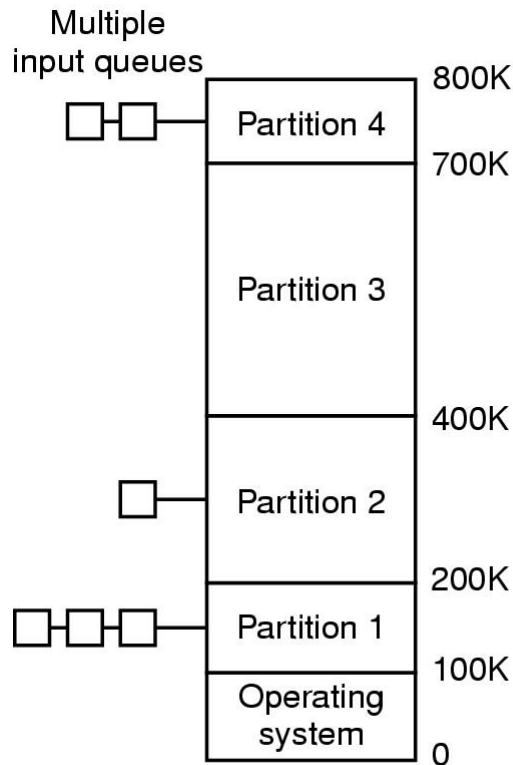
- Fixed memory partitions
  (a) separate input queues of processes for each partition
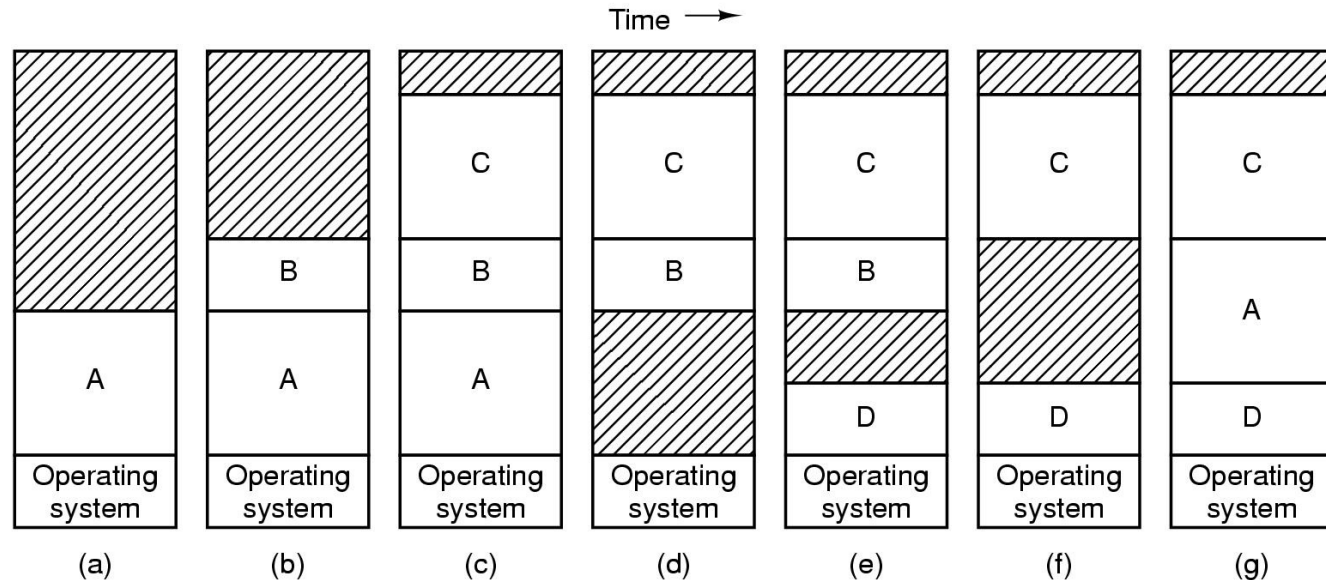 (b) single input queue

Physical Memory addressing

- Consider an instruction that reads from a memory location

    - load reg memory_address

- But programmer doesn't know the memory_address where data will be stored when the process runs!

- Solution: Relocation
    - Programmer assumes a "relative" address, which is converted to a "physical" address by the OS+hardware when the process runs.

Relocation and Protection

- Problem: A programmer doesn't know where a program will be loaded in memory
  - address locations of variables and code routines cannot be absolute
  - must keep a program out of other processes' partitions
- Solution: Use base and limit values
- Relocation
  - Address locations in a program are relative.
  - They are added to a **base value** to map to physical addresses.
- Protection
  - Access to address locations larger than **limit value** results in an error

LIMIT

0

Relative Addresses in original program binary

Physical
MAX

BASE +
LIMIT

BASE

0

Relocated Addresses in Executing Binary

7

What if physical memory is not enough to hold all processes? — Swapping

Time →

| (a) | (b) | (c) | (d) | (e) | (f) | (g) |

(a): A, Operating system
(b): B, A, Operating system
(c): C, B, A, Operating system
(d): C, B, D, Operating system
(e): C, B, D, Operating system
(f): C, D, Operating system
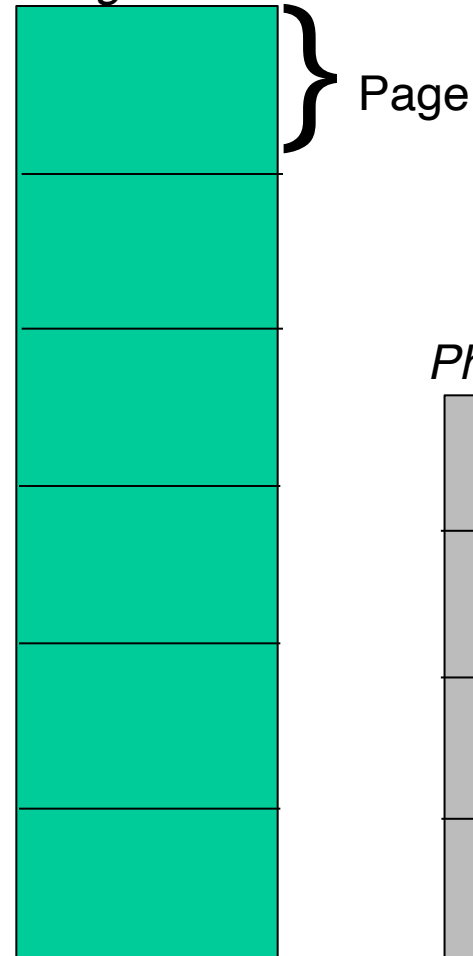(g): C, A, D, Operating system

- Physical memory may not be enough to accommodate the needs of all processes
- Memory allocation changes as
  - processes come into memory
  - leave memory and are *swapped out* to disk
  - Re-enter memory by getting *swapped-in* from disk
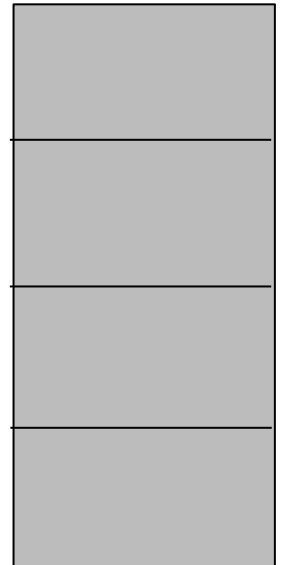- Shaded regions are unused memory

8

Virtual Memory

- Swapping the memory of an entire process is useful when the sum of memory needed by all processes is greater than the total RAM available in the system.

- But sometimes, a single process might require more memory than the total RAM in the system.

- In such cases swapping an entire process is not enough.

- Rather, we need to break up the memory space of a process into smaller equal-sized pieces, called <u>PAGES</u>.

- OS then decides which pages stay in memory and which get moved to disk.

- Virtual memory: means that <u>each process gets an illusion that it has more memory than the physical RAM in the system</u>.

Virtual Address Space of a single Process

} Page

Entire Physical RAM

Memory Management Unit (MMU)

The CPU sends virtual addresses to the MMU

CPU package

CPU

Memory management unit

Memory

Disk controller

Bus

The MMU sends physical addresses to the memory
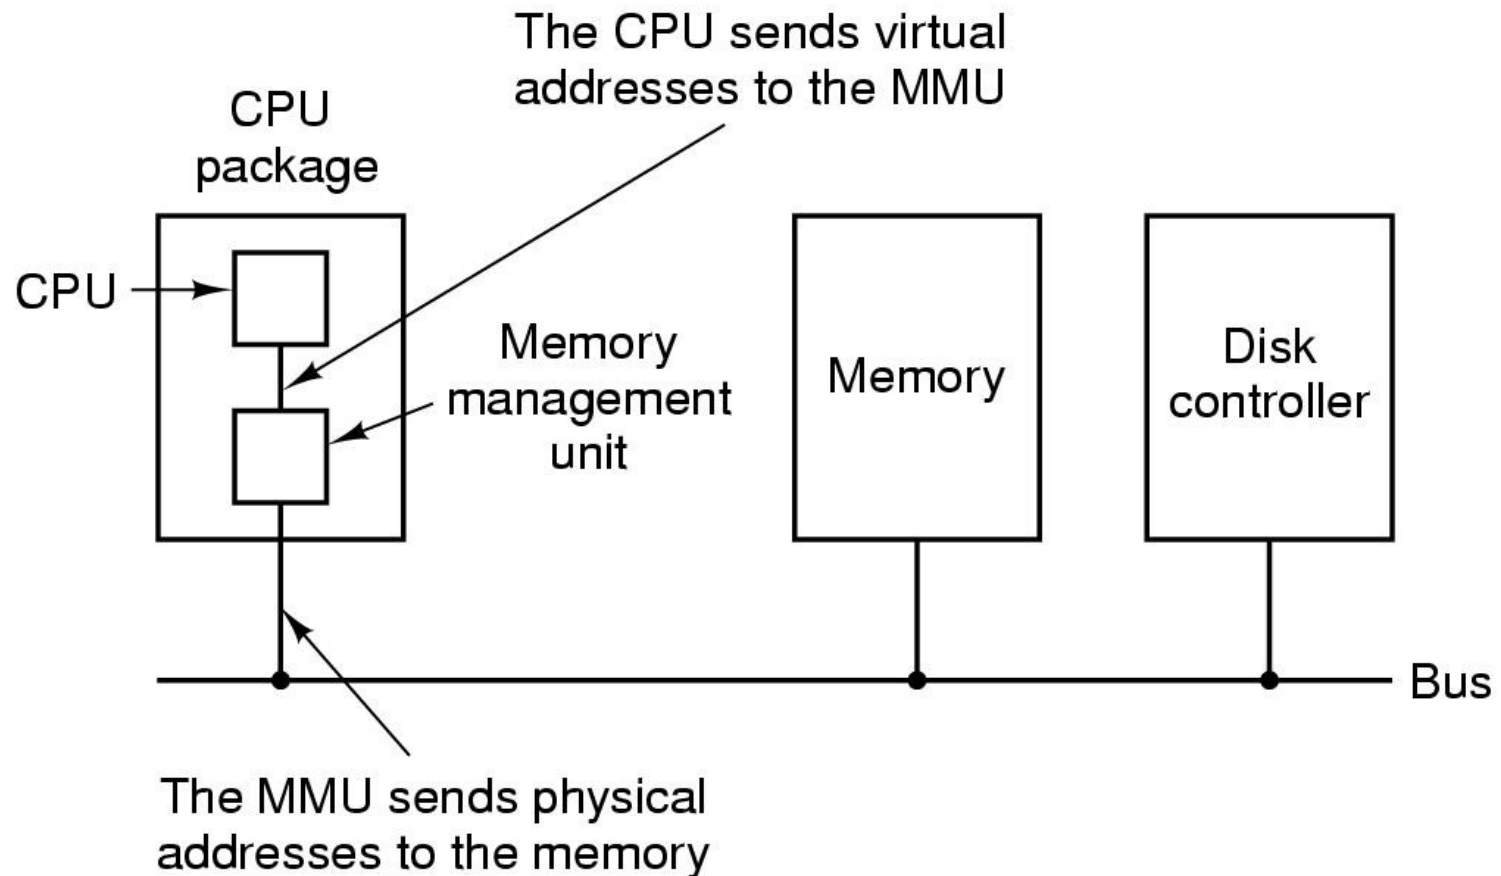
- MMU is a hardware module that accompanies the CPU
- It translates the Virtual Address used by executing instructions to Physical Addresses in the main memory.

Size of address space (in bytes) as a function of address size (in bits)

Number of bits in address	Maximum address space size (bytes)
0	2^0 = 1 byte
1	2^1 = 2 bytes
2	2^2 = 4 bytes
10	2^{10} = 1024 = 1KiB
12	2^{12} = 4KiB
16	2^{16} = 64 KiB
32	2^{32} = 4GiB (Gibibytes)
64	2^{64} = 16 EiB (Exbibytes)

Page Table

- An array that stores the mapping from virtual page numbers to physical numbers

- The OS maintains
  - One page table per userspace process.
  - And usually another page table for kernel memory.

Virtual address space

Virtual address space	
60K-64K	X
56K-60K	X
52K-56K	X
48K-52K	X
44K-48K	7
40K-44K	X
36K-40K	5
32K-36K	X
28K-32K	X
24K-28K	X
20K-24K	3
16K-20K	4
12K-16K	0
8K-12K	6
4K-8K	1
0K-4K	2

} Virtual page

Physical memory address

	28K-32K
	24K-28K
	20K-24K
	16K-20K
	12K-16K
	8K-12K
	4K-8K
	0K-4K

} 0K-4K

Page frame
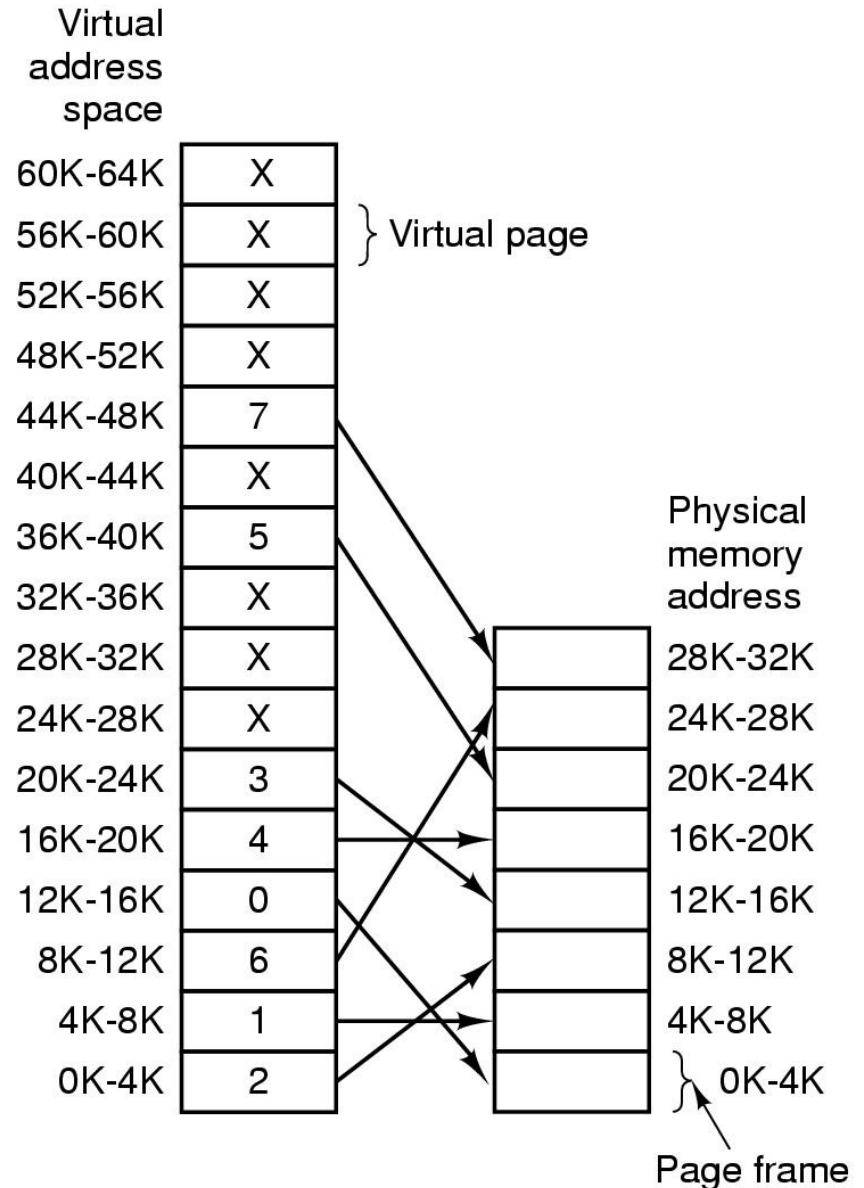
Translating
Virtual address (VA) to physical address (PA)

Virtual Address Space   *Physical RAM*

Byte Address =
Page Number x Page Size +
Byte Offset in the page

VA = VPN x Page Size + Byte Offset

PA = PPN x Page Size + Byte Offset

Byte Offset

6

5

4

3

2

1

0

Byte Offset

Virtual
Page   2
Numbers
(VPN)

Byte Offset

1

0

Physical
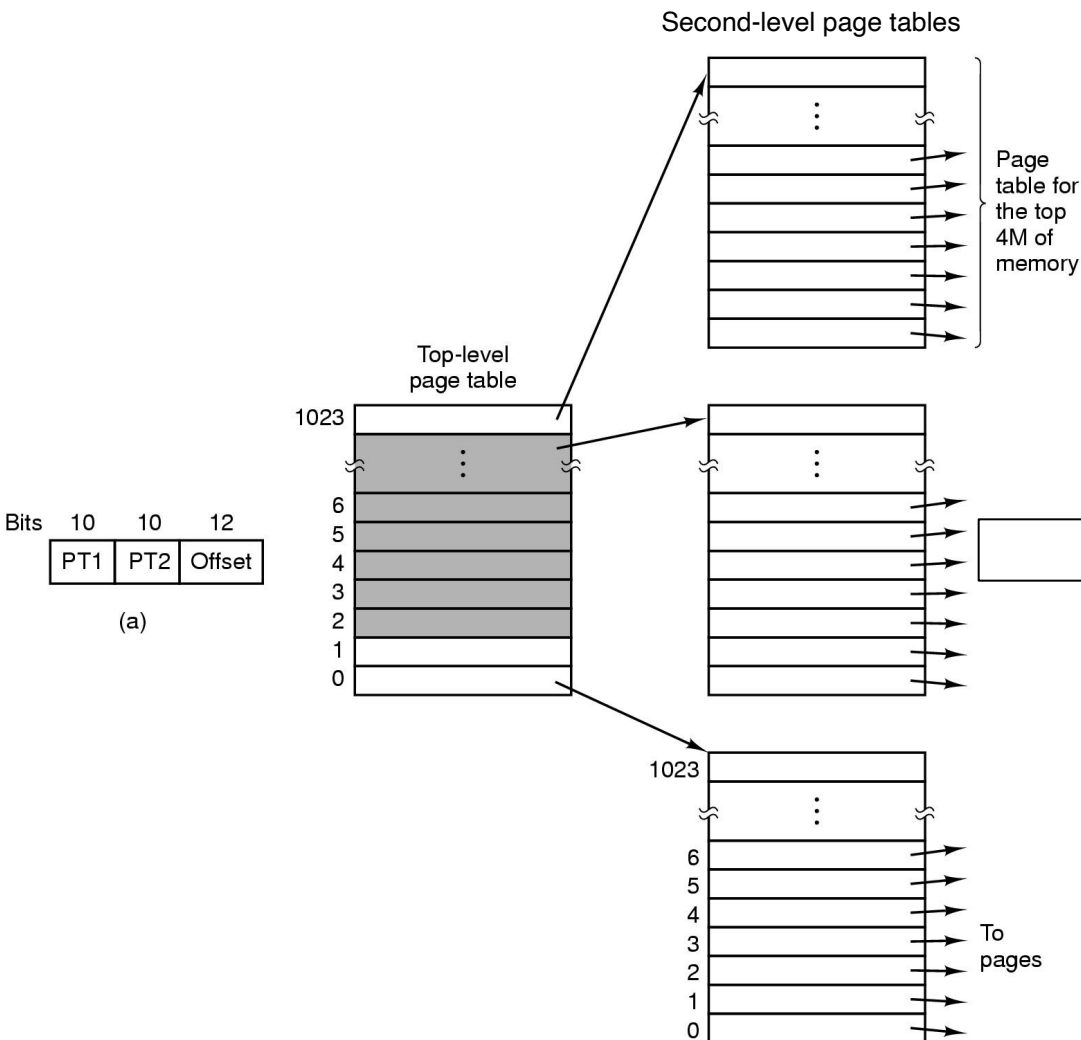Page (frame)
Numbers
(PPN)
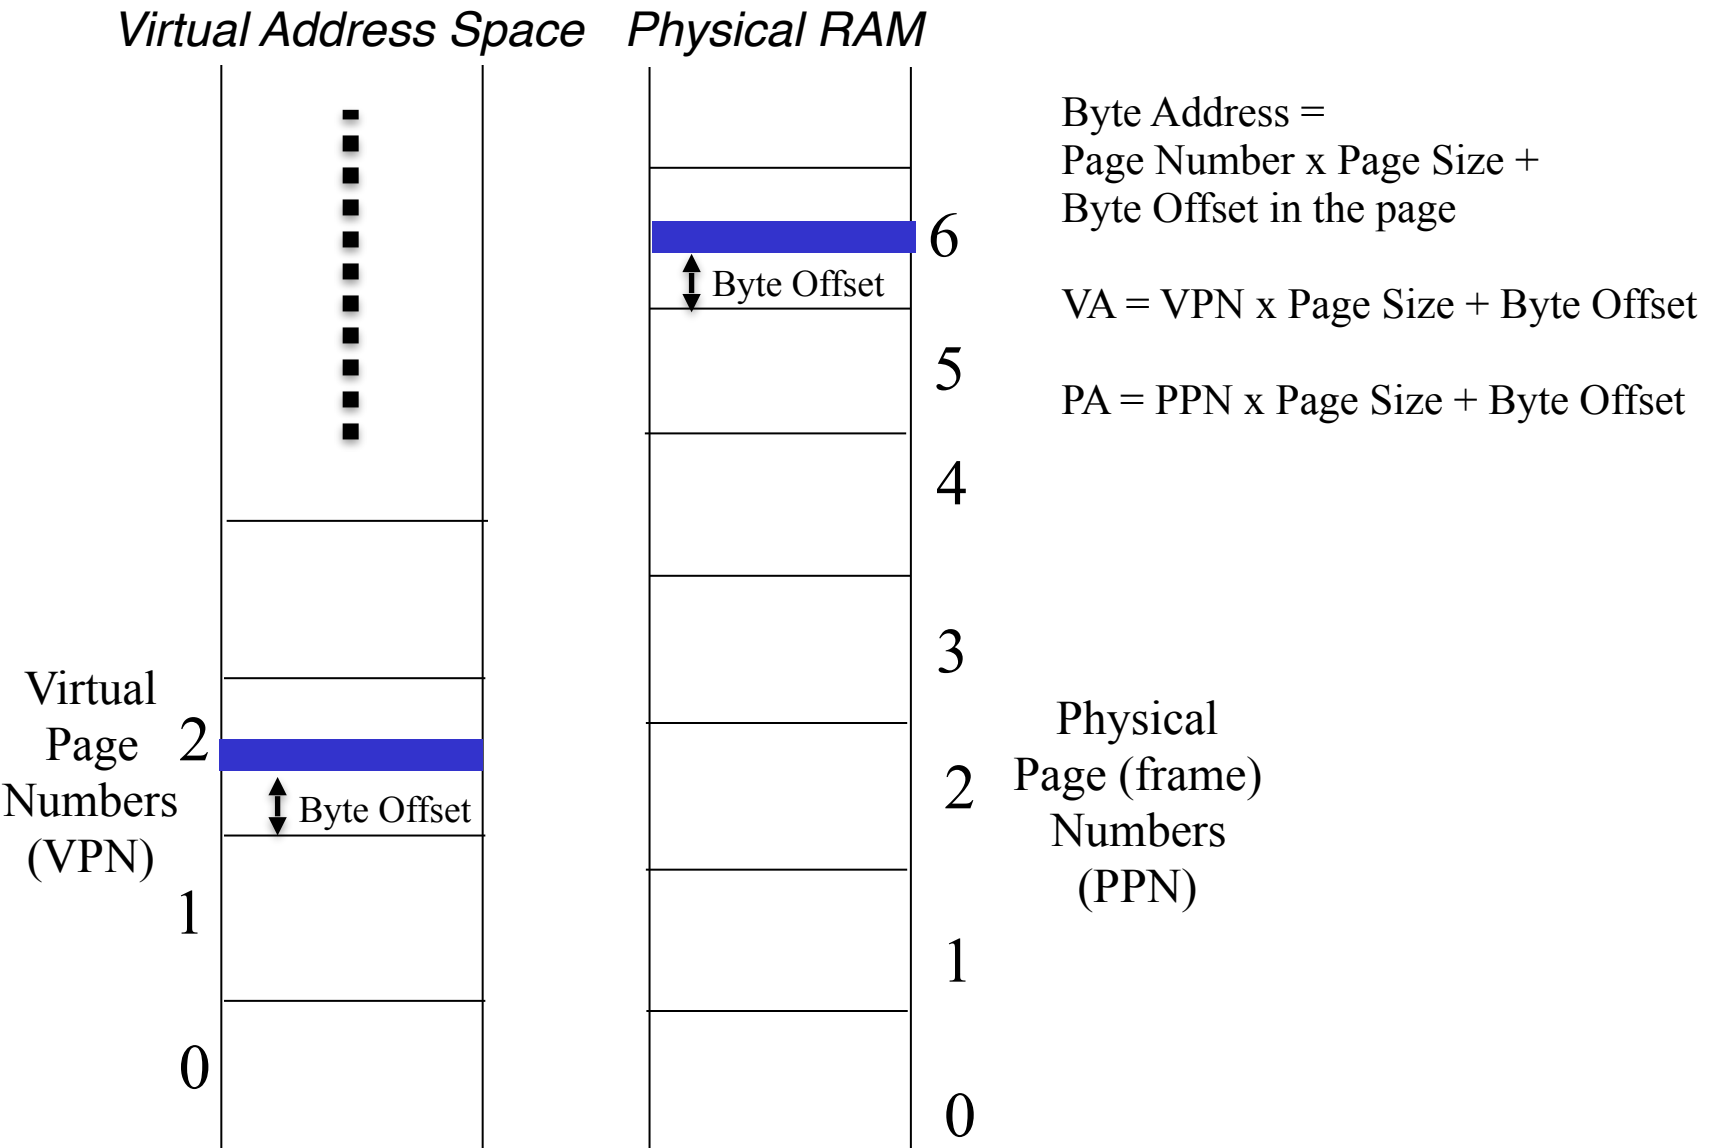
Virtual Address Translation For Small Address Space
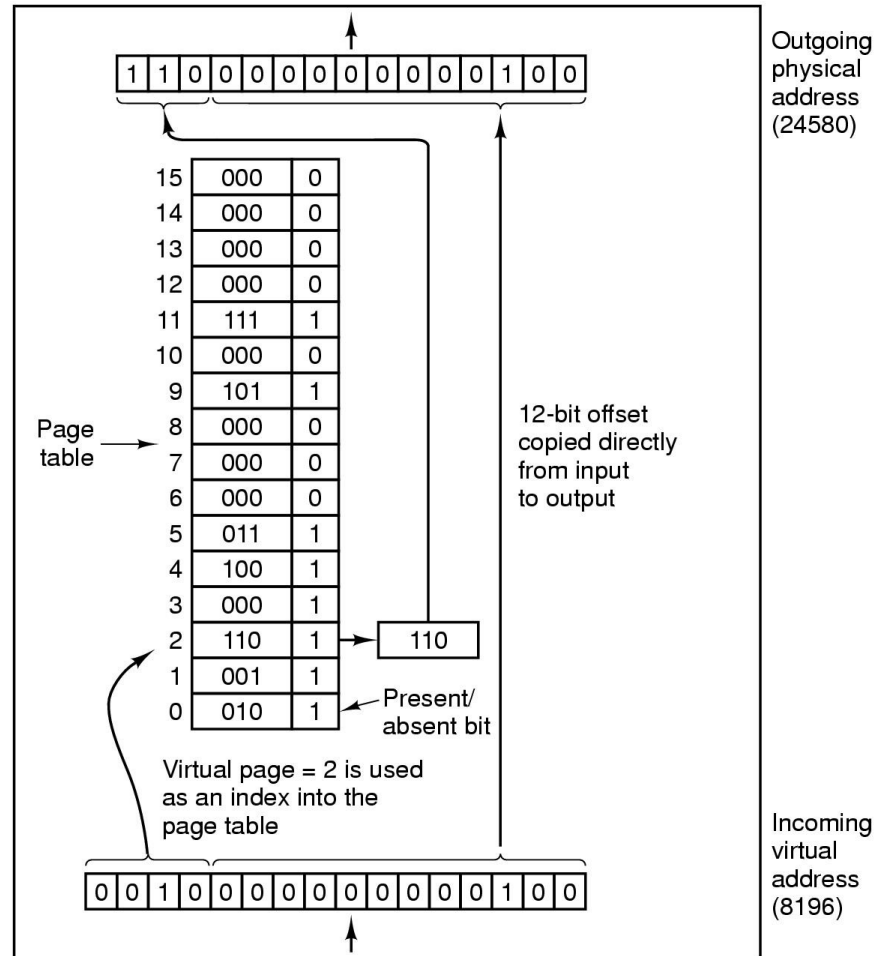
Internal operation of MMU with 16 4 KB pages

Virtual Address Translation For Large Address Space

Second-level page tables

Page table for the top 4M of memory

Top-level page table

Bits 10   10   12
| PT1 | PT2 | Offset |

(a)

1023
6
5
4
3
2
1
0

1023
6
5
4
3
2
1
0

To pages

- 32 bit address with 2 page table fields

- Two-level page tables

- PT too Big for MMU
  - Keep it in main memory

- But how does MMU know where to find PT?
  - Registers (CR2 on Intel)

Typical Page Table Entry (PTE)
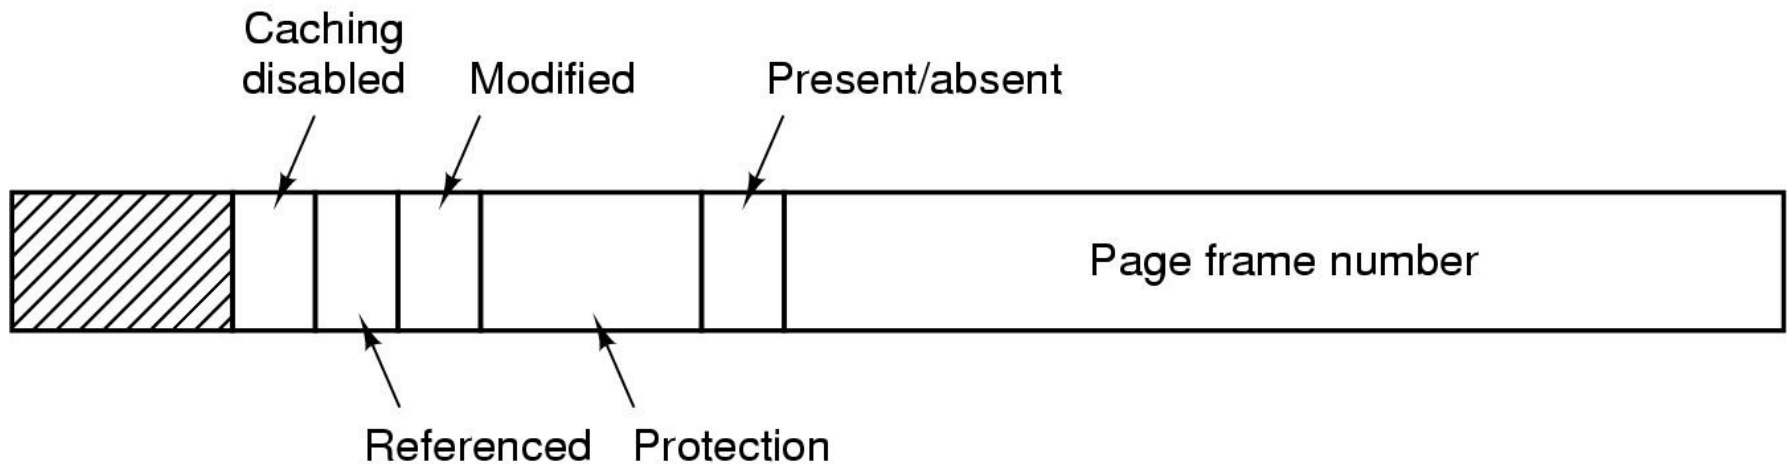
- Page Frame number = physical page number for the virtual page represented by the PTE
- Referenced bit: Whether the page was accessed since last time the bit was reset.
- Modified bit: Also called "Dirty" bit. Whether the page was written to, since the last time the bit was reset.
- Protection bits: Whether the page is readable? writeable? executable? contains higher privilege code/data?
- Present/Absent bit: Whether the PTE contains a valid page frame number. Used for marking swapped/unallocated pages.

TLBs – Translation Lookaside Buffers

Valid	Virtual page	Modified	Protection	Page frame
1	140	1	RW	31
1	20	0	R X	38
1	130	1	RW	29
1	129	1	RW	62
1	19	0	R X	50
1	21	0	R X	45
1	860	1	RW	14
1	861	1	RW	75

- TLB is a small cache that speeds up the translation of virtual addresses to physical addresses.
- TLB is part of the MMU hardware (comes with CPU)
- It is not a Data Cache or Instruction Cache. Those are separate.
- TLB simply caches translations from virtual page number to physical page number so that the MMU don't have to access page-table in memory too often.
- On older x86 processors, TLB had to be "flushed" upon every context switch because there is no field in TLB to identify the process context.
    - Tagged TLB can reduce this overhead

Cold Start Penalty

- Cost of repopulating the TLB (and other caches) upon a context switch.

- Immediately after a context switch, all (or many) of TLB entries are invalidated.

  - On some x86 processors, TLB has to be "flushed" upon every context switch because there is no field in TLB to identify the process context.

- Every memory access by the newly scheduled process may results in a TLB miss.

- MMU must then walk the page-table in main memory to repopulate the missing TLB entry, which takes longer than a cache hit.

Tagged TLB

- A "tag" in each TLB entry identifies the process/ thread context to which the TLB entry belongs

- Thus TLB entries for more than one execution context can be stored simultaneously in the TLB.
  - TLB lookup hardware matches the tag in addition to the virtual page number.

- With tags, context switch no longer requires a complete TLB flush.
  - Reduces cold-start penalty.

Two types of memory translation architectures

❏ Architected Page Tables
- Page table interface defined by ISA and understood by memory translation hardware
- E.g. x86 architecture
- MMU handles TLB miss (in hardware)
- OS handles page faults (in software)
- ISA specifies page table format

❏ Architected TLBs
- TLB interface defined by ISA and understood by MMU
- E.g. alpha architecture
- TLB miss handled by OS (in software)
- ISA does not specify page table format

Impact of Page Size on Page tables

Small page size

- Advantages
  - less internal fragmentation
  - page-in/page-out less expensive

- Disadvantages
  - process that needs more pages has larger page table

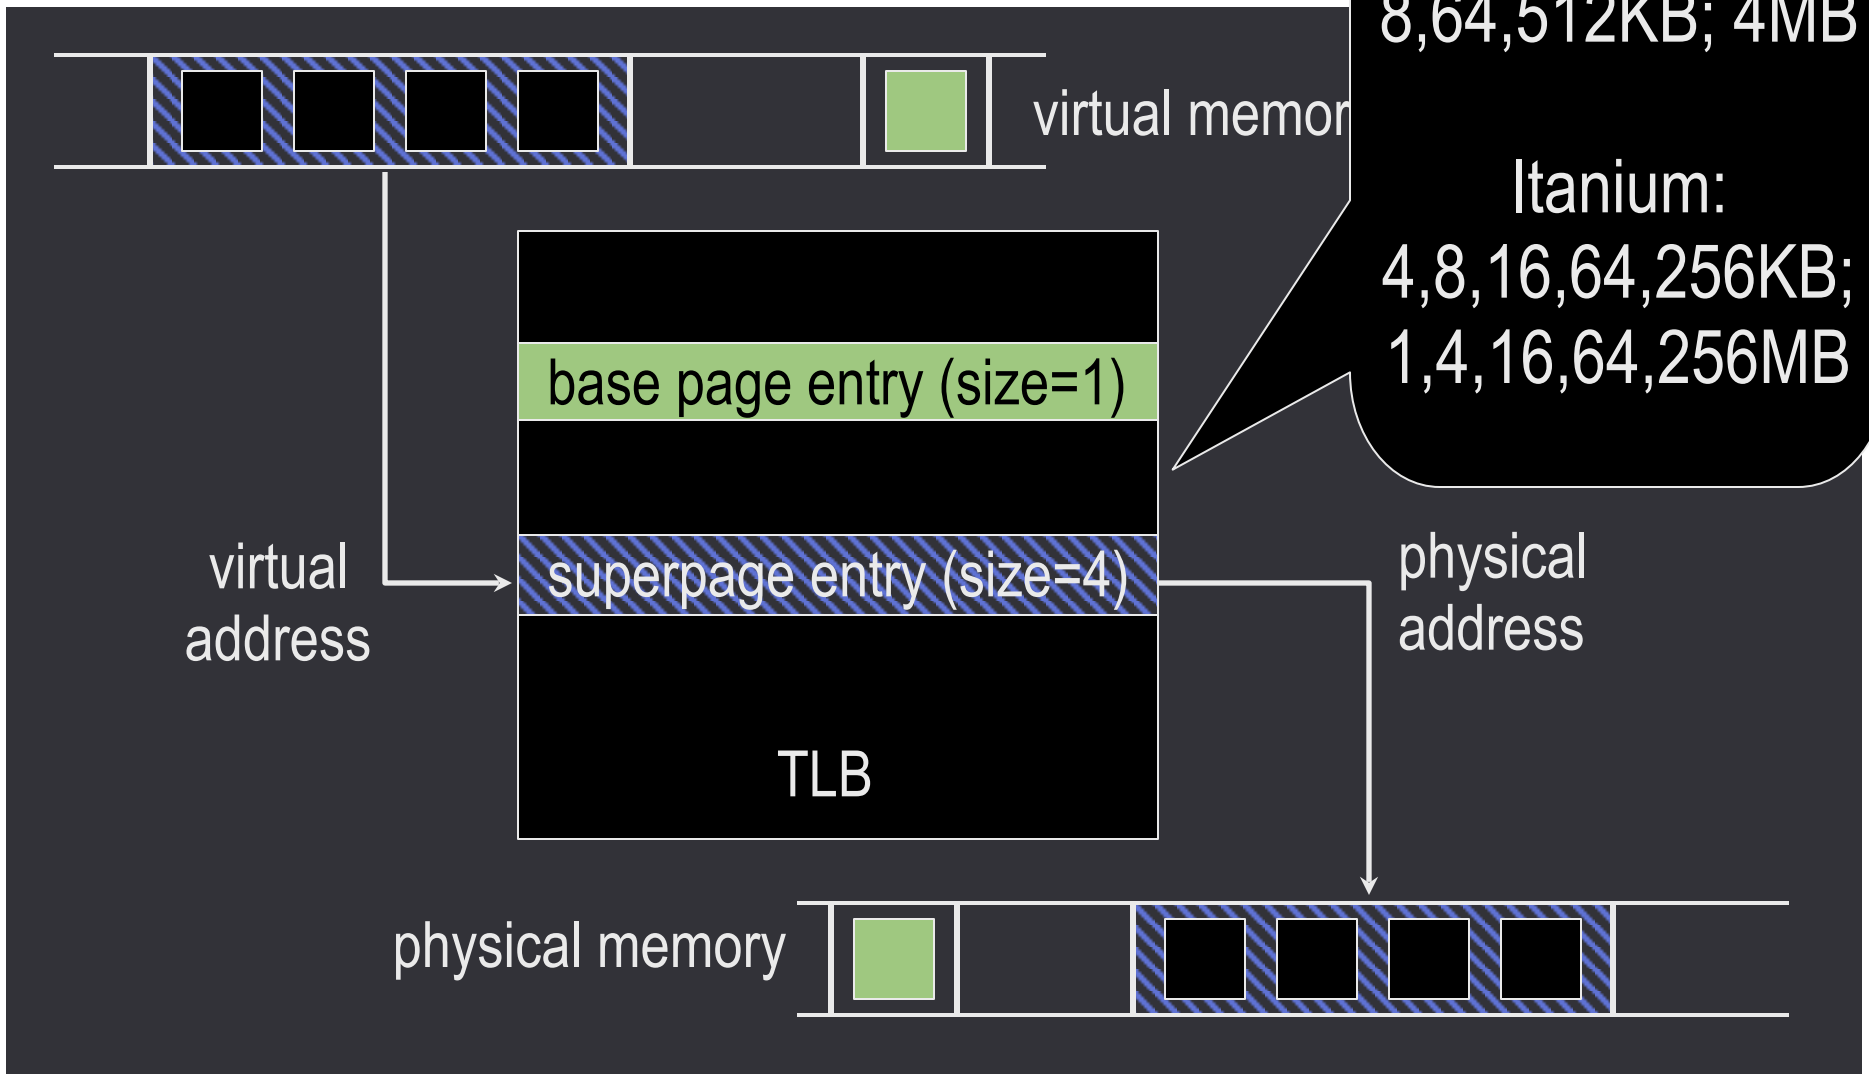- Ref: https://en.wikipedia.org/wiki/Page_(computer_memory)
-

TLB Coverage

- Max amount of memory mapped by TLB
  - Max mount of memory that can be accessed without TLB misses

- TLB Coverage = N x P bytes
  - N = Number of entries in TLB
  - P = Page size in bytes
  - N is fixed by hardware constraints
  - So, to increase TLB Coverage, we must increase P.

- Consider these extreme examples
  - Suppose P = 1 byte
    - TLB Coverage = N bytes only
  - Suppose P = 2^{64} bytes (on a 64-bit ISA)
    - TLB Coverage = N x2^{64}bytes
    - TLB can perform translations for N processes without any TLB misses!

- Of course, both examples above are impractical and meant to illustrate the tradeoffs.

- But what if P is something reasonable, but greater than than the standard 4KB?

- This brings us next to superpages.

Superpages

- Memory pages of larger sizes than standard pages
  - supported by most modern CPUs

- Superpage size = power of 2 x the base page size

- Only one TLB entry per superpage
  - But multiple (identical) page-table entries, one per base page

- Constraints:
  - contiguous (physically and virtually)
  - aligned (physically and virtually)
  - uniform protection attributes
  - one reference bit, one dirty bit

A superpage TLB

virtual memory

base page entry (size=1)

superpage entry (size=4)

virtual address

physical address

TLB

physical memory

Alpha:
8,64,512KB; 4MB

Itanium:
4,8,16,64,256KB;
1,4,16,64,256MB

Quiz

- Consider a machine that has a 32-bit virtual address space and 8KByte page size.

1. What is the total size (in bytes) of the virtual address space for each process?

2. How many bits in a 32-bit address are needed to determine the page number of the address?

3. How many bits in a 32-bit address represent the byte offset into a page?

4. How many page-table entries are present in the page table?

Quiz Answers

- Consider a machine that has a 32-bit virtual address space and 4KByte page size.

1. Total size (in bytes) of the virtual address space for each process = 2^{32} = 4 * 1024 * 1024 *1024 bytes = 4 GB

2. Number of pages in virtual address space = 4GB/8KB = 512*1024 = $2^9*2^{10} = 2^{19}$
   - So the number of bits in a 32-bit address are needed to determine the page number of the address = log2(4GB/8KB) = log2(2^{19}) = 19 bits

3. How many bits in a 32-bit address represent the byte offset into a page?
   - log2(8KB) = log2(2^{13}) = 13
   - Also, 32 – 19 = 13 bits

4. How many page-table entries are present in the page table?
   - Number of PTEs = Number of pages in virtual address = 4GB/8KB = 2^{19} pages

References

- **Chapter 3: Modern Operating Systems, Andrew S. Tanenbaum**

- **X86 architecture**
  http://en.wikipedia.org/wiki/X86

- **Memory segment**
  http://en.wikipedia.org/wiki/Memory_segment

- **Memory model**
  http://en.wikipedia.org/wiki/Memory_model

- **IA-32 Intel Architecture Software Developer's Manual, Volume 1: Basic Architecture**